

Session 9: Advance Hive Assignment 1

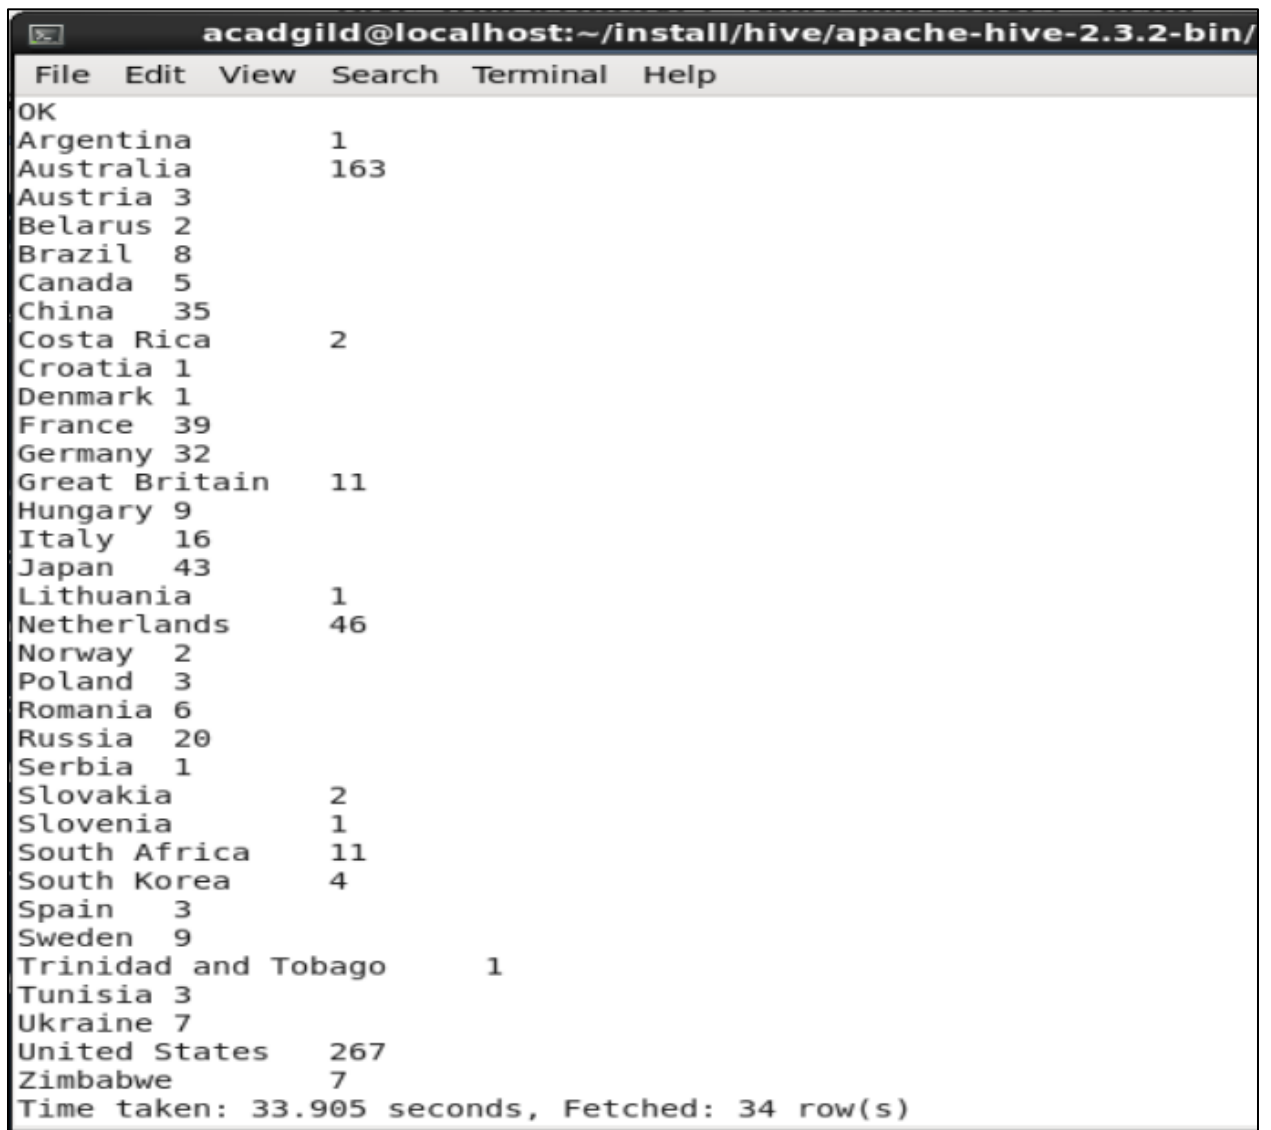
Task 1

1. Write a Hive program to find the number of medals won by each country in swimming.

QUERY:

```
SELECT country,SUM(totalmedal) FROM olympics_data WHERE  
sport='Swimming' GROUP BY country;
```

OUTPUT:



```
acadgild@localhost:~/install/hive/apache-hive-2.3.2-bin/  
File Edit View Search Terminal Help  
OK  
Argentina 1  
Australia 163  
Austria 3  
Belarus 2  
Brazil 8  
Canada 5  
China 35  
Costa Rica 2  
Croatia 1  
Denmark 1  
France 39  
Germany 32  
Great Britain 11  
Hungary 9  
Italy 16  
Japan 43  
Lithuania 1  
Netherlands 46  
Norway 2  
Poland 3  
Romania 6  
Russia 20  
Serbia 1  
Slovakia 2  
Slovenia 1  
South Africa 11  
South Korea 4  
Spain 3  
Sweden 9  
Trinidad and Tobago 1  
Tunisia 3  
Ukraine 7  
United States 267  
Zimbabwe 7  
Time taken: 33.905 seconds, Fetched: 34 row(s)
```

2. Write a Hive program to find the number of medals that India won year wise.

QUERY:

```
SELECT olympicsyear, SUM(totalmedal) FROM olympics_data WHERE country =  
'India' GROUP BY olympicsyear;
```

OUTPUT:

```
Stage: Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.0 sec  
DFS Write: 163 SUCCESS  
Total MapReduce CPU Time Spent: 5 seconds 200 ms  
OK  
2000      1  
2004      1  
2008      3  
2012      6  
Time taken: 31.254 seconds, Fetched: 4 row(s)  
hive>
```

3. Write a Hive Program to find the total number of medals each country won.

QUERY:

```
SELECT country, SUM(totalmedal) FROM olympics_data GROUP BY country;
```

OUTPUT:

Total MapReduce CPU Time Spent: 3 seconds

OK

Afghanistan	2	
Algeria	8	
Argentina	141	
Armenia	10	
Australia	609	
Austria	91	
Azerbaijan	25	
Bahamas	24	
Bahrain	1	
Barbados	1	
Belarus	97	
Belgium	18	
Botswana	1	
Brazil	221	
Bulgaria	41	
Cameroon	20	
Canada	370	
Chile	22	
China	530	
Chinese Taipei	20	
Colombia	13	
Costa Rica	2	
Croatia	81	
Cuba	188	
Cyprus	1	
Czech Republic	81	
Denmark	89	
Dominican Republic		5
Ecuador	1	
Egypt	8	
Eritrea	1	
Estonia	18	
Ethiopia	29	
Finland	118	

France	318	
Gabon	1	
Georgia	23	
Germany	629	
Great Britain	322	
Greece	59	
Grenada	1	
Guatemala	1	
Hong Kong	3	
Hungary	145	
Iceland	15	
India	11	
Indonesia	22	
Iran	24	
Ireland	9	
Israel	4	
Italy	331	
Jamaica	80	
Japan	282	
Kazakhstan	42	
Kenya	39	
Kuwait	2	
Kyrgyzstan	3	
Latvia	17	
Lithuania	30	
Macedonia	1	
Malaysia	3	
Mauritius	1	
Mexico	38	
Moldova	5	
Mongolia	10	
Montenegro	14	
Morocco	11	
Mozambique	1	
Netherlands	318	
New Zealand	52	

```

Nigeria 39
North Korea      21
Norway  192
Panama   1
Paraguay      17
Poland   80
Portugal      9
Puerto Rico   2
Qatar    3
Romania 123
Russia  768
Saudi Arabia   6
Serbia   31
Serbia and Montenegro 38
Singapore      7
Slovakia      35
Slovenia      25
South Africa   25
South Korea   308
Spain   205
Sri Lanka      1
Sudan    1
Sweden  181
Switzerland    93
Syria    1
Tajikistan      3
Thailand    18
Togo      1
Trinidad and Tobago 19
Tunisia  4
Turkey  28
Uganda   1
Ukraine 143
United Arab Emirates 1
United States 1312
Uruguay  1
Uzbekistan      19
Venezuela       4
Vietnam  2
Zimbabwe       7
Time taken: 30.12 seconds, Fetched: 110 row(s)
hive> █

```

4. Write a Hive program to find the number of gold medals each country won.

QUERY:

```
SELECT country, SUM(goldmedal) FROM olympics_data GROUP BY country;
```

OUTPUT:

File	Edit	View	Search	Terminal	Help
OK					
Afghanistan			0		
Algeria	2				
Argentina			49		
Armenia	0				
Australia			163		
Austria	36				
Azerbaijan			6		
Bahamas	11				
Bahrain	0				
Barbados			0		
Belarus	17				
Belgium	2				
Botswana			0		
Brazil	46				
Bulgaria			8		
Cameroon			20		
Canada	168				
Chile	3				
China	234				
Chinese Taipei		2			
Colombia		2			
Costa Rica		0			
Croatia	35				
Cuba	57				
Cyprus	0				
Czech Republic		14			
Denmark	46				
Dominican Republic				3	
Ecuador	0				
Egypt	1				
Eritrea	0				
Estonia	6				
Ethiopia			13		
Finland	11				
France	108				

Gabon	0	
Georgia	6	
Germany	223	
Great Britain	124	
Greece	12	
Grenada	1	
Guatemala	0	
Hong Kong	0	
Hungary	77	
Iceland	0	
India	1	
Indonesia	5	
Iran	10	
Ireland	1	
Israel	1	
Italy	86	
Jamaica	24	
Japan	57	
Kazakhstan	13	
Kenya	11	
Kuwait	0	
Kyrgyzstan	0	
Latvia	3	
Lithuania	5	
Macedonia	0	
Malaysia	0	
Mauritius	0	
Mexico	19	
Moldova	0	
Mongolia	2	
Montenegro	0	
Morocco	2	
Mozambique	1	
Netherlands	101	
New Zealand	18	
Nigeria	6	

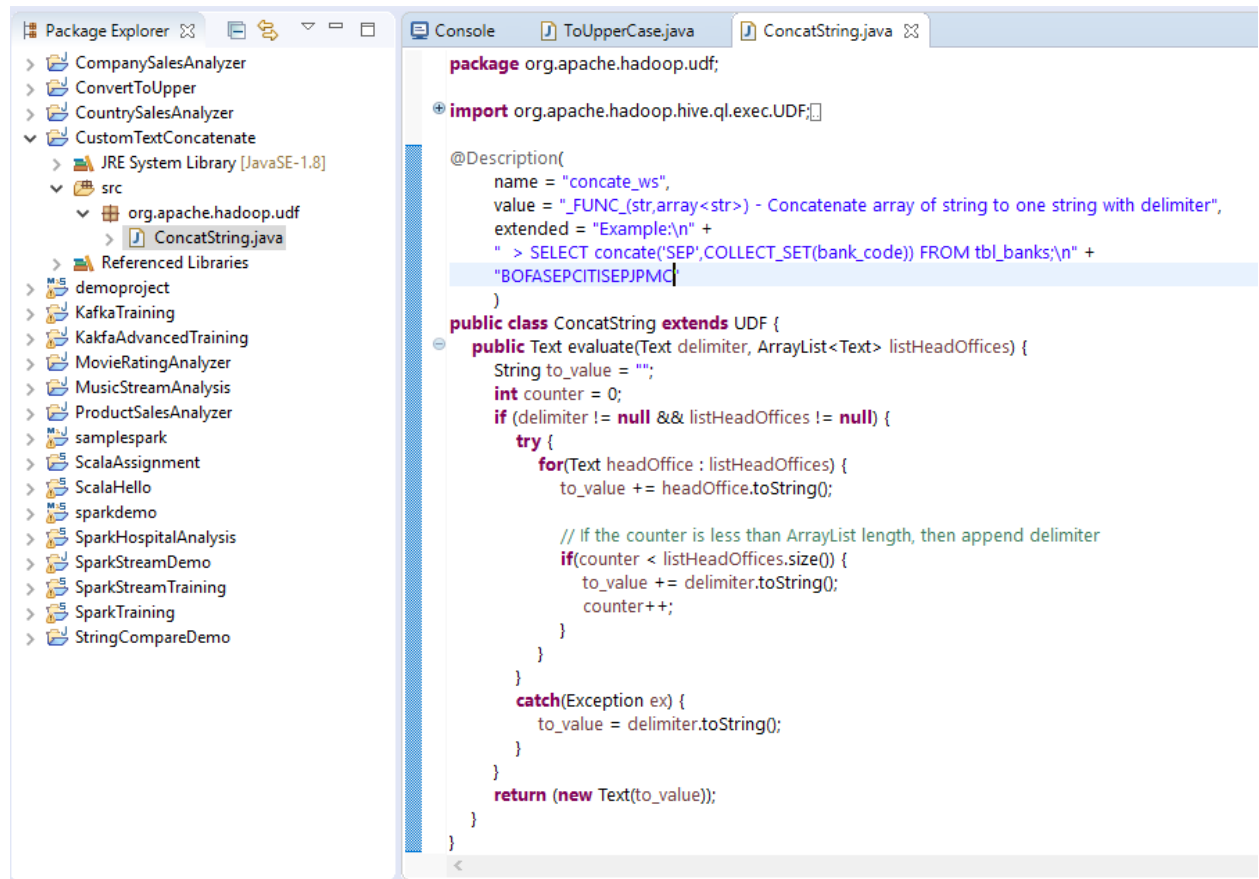
North Korea	6	
Norway	97	
Panama	1	
Paraguay	0	
Poland	20	
Portugal	1	
Puerto Rico	0	
Qatar	0	
Romania	57	
Russia	234	
Saudi Arabia	0	
Serbia	1	
Serbia and Montenegro		11
Singapore	0	
Slovakia	10	
Slovenia	5	
South Africa	10	
South Korea	110	
Spain	19	
Sri Lanka	0	
Sudan	0	
Sweden	57	
Switzerland	21	
Syria	0	
Tajikistan	0	
Thailand	6	
Togo	0	
Trinidad and Tobago		1
Tunisia	2	
Turkey	9	
Uganda	1	
Ukraine	31	
United Arab Emirates		1
United States	552	
Uruguay	0	
Uzbekistan	5	
Venezuela	1	
Vietnam	0	
Zimbabwe	2	

Time taken: 30.528 seconds, Fetched: 110 row(s)
hive> █

Task – 2:

Write a hive UDF that implements functionality of string concat_ws(string SEP, array<string>). This UDF will accept two arguments, one string and one array of string. It will return a single string where all the elements of the array are separated by the SEP.

Program:



```
package org.apache.hadoop.udf;

import org.apache.hadoop.hive.ql.exec.UDF;

@Description(
    name = "concat_ws",
    value = "_FUNC_(str,array<str>) - Concatenate array of string to one string with delimiter",
    extended = "Example:\n" +
        "> SELECT concat('SEP',COLLECT_SET(bank_code)) FROM tbl_banks;\n" +
        "BOFASEP CITISEP JPMC"
)
public class ConcatString extends UDF {
    public Text evaluate(Text delimiter, ArrayList<Text> listHeadOffices) {
        String to_value = "";
        int counter = 0;
        if (delimiter != null && listHeadOffices != null) {
            try {
                for(Text headOffice : listHeadOffices) {
                    to_value += headOffice.toString();

                    // If the counter is less than ArrayList length, then append delimiter
                    if(counter < listHeadOffices.size()) {
                        to_value += delimiter.toString();
                        counter++;
                    }
                }
            } catch (Exception ex) {
                to_value = delimiter.toString();
            }
        }
        return (new Text(to_value));
    }
}
```

Build and export as jar. Then register the jar in Hive as temporary function,

```
text.jar does not exist
hive> add jar /home/acadgild/workspace/Hive/jars/ConcatText.jar;
Added [/home/acadgild/workspace/Hive/jars/ConcatText.jar] to class path
Added resources: [/home/acadgild/workspace/Hive/jars/ConcatText.jar]
hive> create temporary function concat_ws as 'org.apache.hadoop.udf.ConcatString';
OK
Time taken: 0.052 seconds
hive> list jars;
/home/acadgild/workspace/Hive/jars/ConcatText.jar
hive> SELECT concat_ws('SEP',COLLECT_SET(bank_code)) FROM tbl_banks;
```

Execution:

```
hive> list jars;
/home/acadgild/workspace/Hive/jars/ConcatText.jar
hive> SELECT concat_ws('SEP',COLLECT SET(bank code)) FROM tbl_banks;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190318220200_2e758f92-4640-4c46-a6d8-f50783f92ca1
Total jobs = 1
Launching Job 1 out of 1
```

The above command intends to concatenate bank code with delimiter as "SEP"

```
hive> select * from tbl_banks;
OK
101      Citibank      CITI      New York
102      Bank of America BOFA      New York
103      JP Morgan Chase JPMC      Baltimore
Time taken: 0.431 seconds, Fetched: 3 row(s)
hive>
```

Final result will be as follows,

```
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 1 Cumulative CPU: 15.93 sec HDFS Read: 33131
HDFS Write: 118 SUCCESS
Total MapReduce CPU Time Spent: 15 seconds 930 msec
OK
BOFASEPJCITISEPJPMC
Time taken: 87.27 seconds, Fetched: 1 row(s)
hive>
Create temporary function toUpperCase as org.apache.hadoop.hive.ql.udf.UDFUpperCase;
```

Task – 3:

Link: <https://acadgild.com/blog/transactions-in-hive/>

Refer the above given link for transactions in Hive and implement the operations given in the blog using your own sample data set and send us the screenshot.

Execution:

```
ns. Consider using a different execution engine (i.e. spark, tez) or using Hive
1.X releases.
hive> show databases
> ;
OK
custom
default
Time taken: 8.991 seconds, Fetched: 2 row(s)
hive> create database acadgild;
OK
Time taken: 0.37 seconds
hive> show databases;
OK
acadgild
custom
default
Time taken: 0.063 seconds, Fetched: 3 row(s)
hive> use acadgild;
OK
Time taken: 0.054 seconds
hive> show tables;
OK
Time taken: 0.073 seconds
hive> █
```

Properties to set appropriately in Hive shell to work with hive transactions in table,

```
Time taken: 0.073 seconds
hive> set hive.support.concurrency=true;
hive> set hive.enforce.bucketing=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> set hive.txn.manager=org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
hive> set hive.compactor.initiator.on=true;
hive> set hive.compactor.worker.threads=2;
hive> █
```

CREATE NEW TABLE:

```
CREATE TABLE tbl_banks(bank_id int,bank_name string,bank_code string,ho_location string)
clustered by (bank_id) into 5 buckets stored as orc TBLPROPERTIES('transactional'='true');
```

```

hive> set hive.compactor.worker.threads=2;
hive> CREATE TABLE tbl_banks(bank_id int,bank_name string,bank_code string,ho_lo
cation string) clustered by (bank_id) into 5 buckets stored as orc TBLPROPERTIES
('transactional'='true');
OK
Time taken: 1.469 seconds
hive> show tables;
OK
tbl_banks
Time taken: 0.085 seconds, Fetched: 1 row(s)
hive>

```

Insert data in to Hive table,

```

INSERT INTO TABLE tbl_banks VALUES (101,'Citibank','CITI','New York'),(102,'Bank of
America','BOFA','New York'),(103,'JP Morgan Chase','JPMC','Washington'),(104,'Morgan
Stanley','MRST','Los Angeles');

```

```

tbl_banks
Time taken: 0.085 seconds, Fetched: 1 row(s)
hive> INSERT INTO TABLE tbl_banks VALUES (101,'Citibank','CITI','New York'),(102
,'Bank of America','BOFA','New York'),(103,'JP Morgan Chase','JPMC','Washington'
),(104,'Morgan Stanley','MRST','Los Angeles');
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the futu
re versions. Consider using a different execution engine (i.e. spark, tez) or us
ing Hive 1.X releases.
Query ID = acadgild_20190317001942_5f35648a-a4e5-4f66-b49e-2c436922cbcb9
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1552760528223_0001, Tracking URL = http://localhost:8088/prox

```

```

2019-03-17 00:21:04,477 Stage-1 map = 100%,  reduce = 87%, Cumulative CPU 16.52
sec
2019-03-17 00:21:05,721 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 20.47
sec
MapReduce Total cumulative CPU time: 20 seconds 470 msec
Ended Job = job_1552760528223_0001
Loading data to table acadgild.tbl_banks
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5   Cumulative CPU: 20.47 sec   HDFS Read: 29090
HDFS Write: 4109 SUCCESS
Total MapReduce CPU Time Spent: 20 seconds 470 msec
OK
Time taken: 86.843 seconds
hive>

```

Query the inserted data and display using SELECT statement,

```

SELECT * FROM tbl_banks;

```

```

OK
Time taken: 86.843 seconds
hive> SELECT * FROM tbl_banks;
OK
101      Citibank          CITI      New York
102      Bank of America BOFA      New York
103      JP Morgan Chase JPMC      Washington
104      Morgan Stanley MRST      Los Angeles
Time taken: 0.427 seconds, Fetched: 4 row(s)
hive>

```

Try to update a bucketed column. In tbl_banks table, bank_id column is classified as bucketed. So, it should not be allowed to update and show error message as below,

```
Time taken: 00.043 seconds
hive> SELECT * FROM tbl_banks;
OK
101      Citibank      CITI      New York
102      Bank of America BOFA      New York
103      JP Morgan Chase JPMC      Washington
104      Morgan Stanley MRST      Los Angeles
Time taken: 0.427 seconds, Fetched: 4 row(s)
hive> UPDATE tbl_banks SET bank_id=106 WHERE bank_id=103;
FAILED: SemanticException [Error 10302]: Updating values of bucketing columns is
not supported. Column bank_id.
hive>
```

Update a non-bucketed column value and the result will appear below,

UPDATE tbl_banks SET ho_location='Baltimore' WHERE bank_id=103;

```
not supported. Column bank_id.
hive> UPDATE tbl_banks SET ho_location='Baltimore' WHERE bank_id=103;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future
versions. Consider using a different execution engine (i.e. spark, tez) or using
Hive 1.X releases.
Query ID = acadgild_20190317003502_d9ada84c-1409-4edb-9d34-de6461e13377
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1552760528223_0002, Tracking URL = http://localhost:8088/proxy/application_1552760528223_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1552760528223_0002
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2019-03-17 00:35:14,750 Stage-1 map = 0%, reduce = 0%
```

```
sec
2019-03-17 00:36:39,552 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 26.78
sec
MapReduce Total cumulative CPU time: 26 seconds 780 msec
Ended Job = job_1552760528223_0002
Loading data to table acadgild.tbl_banks
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 26.78 sec HDFS Read: 55269
HDFS Write: 1164 SUCCESS
Total MapReduce CPU Time Spent: 26 seconds 780 msec
OK
Time taken: 99.138 seconds
hive>
```

```

Time taken: 99.138 seconds
hive> SELECT * FROM tbl_banks;
OK
101      Citibank          CITI      New York
102      Bank of America BOFA      New York
103      JP Morgan Chase JPMC      Baltimore
104      Morgan Stanley MRST      Los Angeles
Time taken: 0.412 seconds, Fetched: 4 row(s)
hive>

```

Delete an existing row in hive table,

DELETE FROM tbl_banks WHERE bank_id=104;

```

Time taken: 0.412 seconds, Fetched: 4 row(s)
hive> DELETE FROM tbl_banks WHERE bank_id=104;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190317004005_90d33068-d340-49ec-aa2c-f680c42c565f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1552760528223_0003, Tracking URL = http://localhost:8088/proxy/application_1552760528223_0003/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1552760528223_0003

MapReduce Total cumulative CPU time: 25 seconds 140 msec
Ended Job = job_1552760528223_0003
Loading data to table acadgild.tbl_banks
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 25.14 sec HDFS Read: 51995
HDFS Write: 777 SUCCESS
Total MapReduce CPU Time Spent: 25 seconds 140 msec
OK
Time taken: 95.39 seconds
hive>

```

Result:

```

Time taken: 95.39 seconds
hive> SELECT * FROM tbl_banks;
OK
101      Citibank          CITI      New York
102      Bank of America BOFA      New York
103      JP Morgan Chase JPMC      Baltimore
Time taken: 0.308 seconds, Fetched: 3 row(s)
hive>

```